

```
In [1]: # 10/27/2025  
# BSAN 360 Lab 4  
# Lucas Doyle
```

```
In [2]: # 1. Download and read your data into a data frame. Check that the process does not  
import pandas as pd
```

```
In [3]: # 1  
import os  
os.chdir("C:/Users/lucas/Programming")  
print(os.getcwd())
```

C:\Users\lucas\Programming

```
In [4]: # 1  
sleep = pd.read_csv('./Sleep_health_and_lifestyle_dataset.csv')  
sleep.head()
```

Out[4]:

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Pi
0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight	
1	2	Male	28	Doctor	6.2	6	60	8	Normal	
2	3	Male	28	Doctor	6.2	6	60	8	Normal	
3	4	Male	28	Sales Representative	5.9	4	30	8	Obese	
4	5	Male	28	Sales Representative	5.9	4	30	8	Obese	



```
In [5]: # 2. How many rows and columns does your data frame have? Check that this is the co  
print("Shape of dataset:", sleep.shape)  
print("Number of rows:", sleep.shape[0])  
print("Number of columns:", sleep.shape[1])
```

Shape of dataset: (374, 13)

Number of rows: 374

Number of columns: 13

```
In [6]: # 2  
# 372 rows  
# 13 columns
```

```
In [7]: # 3 Check the first and last few rows of data. Does everything look correct?  
sleep.head()
```

Out[7]:

Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Pi
0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight
1	2	Male	28	Doctor	6.2	6	60	8	Normal
2	3	Male	28	Doctor	6.2	6	60	8	Normal
3	4	Male	28	Sales Representative	5.9	4	30	8	Obese
4	5	Male	28	Sales Representative	5.9	4	30	8	Obese



In [8]:

```
# 3
sleep.tail()
```

Out[8]:

Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Pi
369	370	Female	59	Nurse	8.1	9	75	3	Overweight
370	371	Female	59	Nurse	8.0	9	75	3	Overweight
371	372	Female	59	Nurse	8.1	9	75	3	Overweight
372	373	Female	59	Nurse	8.1	9	75	3	Overweight
373	374	Female	59	Nurse	8.1	9	75	3	Overweight



In [9]:

```
# 4. Make sure that column and row labels got processed correctly. Re-Label columns
sleep.columns
```

Out[9]:

```
Index(['Person ID', 'Gender', 'Age', 'Occupation', 'Sleep Duration',
       'Quality of Sleep', 'Physical Activity Level', 'Stress Level',
       'BMI Category', 'Blood Pressure', 'Heart Rate', 'Daily Steps',
       'Sleep Disorder'],
      dtype='object')
```

In [10]:

```
# 4
sleep.rename(columns={'Physical Activity Level': 'Activity'}, inplace=True)
sleep.columns
```

```
Out[10]: Index(['Person ID', 'Gender', 'Age', 'Occupation', 'Sleep Duration',  
                 'Quality of Sleep', 'Activity', 'Stress Level', 'BMI Category',  
                 'Blood Pressure', 'Heart Rate', 'Daily Steps', 'Sleep Disorder'],  
                dtype='object')
```

```
In [11]: # 5. Look at the descriptive statistics and perform the following checks and any ne  
sleep.describe()
```

```
Out[11]:
```

	Person ID	Age	Sleep Duration	Quality of Sleep	Activity	Stress Level	Heart Rate
count	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000
mean	187.500000	42.184492	7.132086	7.312834	59.171123	5.385027	70.165775
std	108.108742	8.673133	0.795657	1.196956	20.830804	1.774526	4.135676
min	1.000000	27.000000	5.800000	4.000000	30.000000	3.000000	65.000000
25%	94.250000	35.250000	6.400000	6.000000	45.000000	4.000000	68.000000
50%	187.500000	43.000000	7.200000	7.000000	60.000000	5.000000	70.000000
75%	280.750000	50.000000	7.800000	8.000000	75.000000	7.000000	72.000000
max	374.000000	59.000000	8.500000	9.000000	90.000000	8.000000	86.000000



```
In [12]: # 5  
sleep.describe(include='object')
```

```
Out[12]:
```

	Gender	Occupation	BMI Category	Blood Pressure	Sleep Disorder
count	374	374	374	374	155
unique	2	11	4	25	2
top	Male	Nurse	Normal	130/85	Sleep Apnea
freq	189	73	195	99	78

```
In [13]: # 5  
# a What information is provided in this column?  
# Person ID - An identifier for each person  
# Age - The participants range from 27 to 59. (mean is 42)  
# Sleep Duration - Hours of sleep per night (min 5.8, max 8.5, mean ~ 7.1). Reflects variety.  
# Quality of sleep - Self-rated quality on a 1-10 scale (range 4-9, mean ~ 7.3). Involves self-assessment.  
# Activity - Daily physical-activity score (30-90, mean ~ 59). Suggests moderate to low activity levels.  
# Stress Level - Self rated stress on a 1-10 scale (3-8, mean ~ 5.4). Indicates moderate stress levels.  
# Heart rate - Average resting heart rate in bpm (65-86, mean ~ 70). ALL values within normal range.  
# Daily Steps - Steps taken per day (3 000-10 000, mean ~ 6 817). Represents mostly sedentary individuals.  
# gender - Two categories - Male (189) and Female (185). Balanced sample; no missing data.  
# occupation - Eleven job types most frequent = Nurse (73). Reflects varied professional backgrounds.  
# BMI Category - Four categories - Normal (most common 195), Overweight, Obese, Underweight.  
# Blood Pressure - 25 unique readings (most common 130/85). Indicates variety in cardiovascular health.  
# Sleep Disorder - Two diagnoses (Sleep Apnea = 78, Insomnia = 77) with 155 non-null entries.
```

```
In [14]: # 5 b What type of data is given in this column? The types we've seen most are Nume  
# (integer or real number), Character Strings, Boolean, Datetime. If the data type  
# correct, correct it as best you can. (A full resetting of data types may require  
# some handling of missing values, which we will cover on Wednesday, so no worries  
sleep.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 374 entries, 0 to 373  
Data columns (total 13 columns):  
 #   Column           Non-Null Count  Dtype     
---  --  
 0   Person ID        374 non-null    int64    
 1   Gender           374 non-null    object    
 2   Age              374 non-null    int64    
 3   Occupation       374 non-null    object    
 4   Sleep Duration   374 non-null    float64   
 5   Quality of Sleep 374 non-null    int64    
 6   Activity          374 non-null    int64    
 7   Stress Level     374 non-null    int64    
 8   BMI Category     374 non-null    object    
 9   Blood Pressure   374 non-null    object    
 10  Heart Rate        374 non-null    int64    
 11  Daily Steps       374 non-null    int64    
 12  Sleep Disorder   155 non-null    object    
dtypes: float64(1), int64(7), object(5)  
memory usage: 38.1+ KB
```

```
In [15]: # 5 C What is the range of values in this column? Check the min and max values to m  
sleep.describe()
```

```
Out[15]:
```

	Person ID	Age	Sleep Duration	Quality of Sleep	Activity	Stress Level	Heart Rate
count	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000
mean	187.500000	42.184492	7.132086	7.312834	59.171123	5.385027	70.165775
std	108.108742	8.673133	0.795657	1.196956	20.830804	1.774526	4.135676
min	1.000000	27.000000	5.800000	4.000000	30.000000	3.000000	65.000000
25%	94.250000	35.250000	6.400000	6.000000	45.000000	4.000000	68.000000
50%	187.500000	43.000000	7.200000	7.000000	60.000000	5.000000	70.000000
75%	280.750000	50.000000	7.800000	8.000000	75.000000	7.000000	72.000000
max	374.000000	59.000000	8.500000	9.000000	90.000000	8.000000	86.000000

```
In [16]: # 5 d Are there any missing values in this column? If so, how do you plan to handle  
sleep.isnull().sum()  
# Only the Sleep Disorder column has missing values (219).
```

```
Out[16]: Person ID      0  
Gender          0  
Age            0  
Occupation     0  
Sleep Duration  0  
Quality of Sleep 0  
Activity        0  
Stress Level    0  
BMI Category    0  
Blood Pressure   0  
Heart Rate       0  
Daily Steps      0  
Sleep Disorder    219  
dtype: int64
```

```
In [17]: # 6. Next, we will focus on the project questions and figure out a roadmap for your  
# your project questions from your assignment here, and for each question, determin  
  
# Project Question:  
# How do lifestyle factors like stress and activity relate to sleep duration and sl  
  
# Columns needed:  
# - Sleep Duration  
# - Quality of Sleep  
# - Stress Level  
# - Activity  
  
# Plan:  
# Compare sleep duration and sleep quality against stress and activity levels  
# to see if healthier daytime behaviors lead to better sleep.
```

```
In [18]: # 7. Create a new data frame that contains all the columns you'll need for your pro  
# computer's memory allows, create individual data frames for each of your research  
  
sleep_project = sleep[['Sleep Duration', 'Quality of Sleep', 'Stress Level', 'Activ  
  
sleep_project.head()
```

```
Out[18]:  


|   | Sleep Duration | Quality of Sleep | Stress Level | Activity |
|---|----------------|------------------|--------------|----------|
| 0 | 6.1            | 6                | 6            | 42       |
| 1 | 6.2            | 6                | 8            | 60       |
| 2 | 6.2            | 6                | 8            | 60       |
| 3 | 5.9            | 4                | 8            | 30       |
| 4 | 5.9            | 4                | 8            | 30       |


```

```
In [19]: # 7  
stress_sleep = sleep[['Stress Level', 'Sleep Duration', 'Quality of Sleep']]  
activity_sleep = sleep[['Activity', 'Sleep Duration', 'Quality of Sleep']]
```

```
In [20]: # 8 Going back to your research questions themselves, determine what type of statis
# you'll need to answer each one. Some examples could be ANOVA to determine if ther
# differences among groups, regression for prediction, time series analysis for for
# classification or clustering methods
```

```
# Analyses I Will Use:
# - Correlation analysis to see if there is a relationship between stress, activit
# - Simple linear regression to predict sleep duration or sleep quality from one f
# - Multiple linear regression to test how both stress and activity together affect
```

```
In [21]: # Project assignment 2
# PROJECT DATA CLEANING AND PREPARATION
sleep_project = sleep[['Sleep Duration', 'Quality of Sleep', 'Stress Level', 'Activ
sleep_project.head()
```

	Sleep Duration	Quality of Sleep	Stress Level	Activity
0	6.1	6	6	42
1	6.2	6	8	60
2	6.2	6	8	60
3	5.9	4	8	30
4	5.9	4	8	30

```
In [22]: # Looking for missing values
sleep_project.isnull().sum()
```

```
Out[22]: Sleep Duration      0
Quality of Sleep      0
Stress Level        0
Activity            0
dtype: int64
```

```
In [23]: # All four columns (Sleep Duration, Quality of Sleep, Stress Level, Activity) have
# No further handling of missing data is required.
# This confirms the dataset is complete and ready for analysis.
```

```
In [24]: # Data transformations
# making sure everything is numeric
sleep_project.dtypes
```

```
Out[24]: Sleep Duration      float64
Quality of Sleep      int64
Stress Level        int64
Activity            int64
dtype: object
```

```
In [25]: # All columns are numeric (float or int).
# No type conversions are needed.
```

```
In [26]: # Final verification of cleaned data
#isplay descriptive statistics to confirm all values look reasonable
```

```
sleep_project.describe()
```

Out[26]:

	Sleep Duration	Quality of Sleep	Stress Level	Activity
count	374.000000	374.000000	374.000000	374.000000
mean	7.132086	7.312834	5.385027	59.171123
std	0.795657	1.196956	1.774526	20.830804
min	5.800000	4.000000	3.000000	30.000000
25%	6.400000	6.000000	4.000000	45.000000
50%	7.200000	7.000000	5.000000	60.000000
75%	7.800000	8.000000	7.000000	75.000000
max	8.500000	9.000000	8.000000	90.000000

In [27]:

```
# The dataset has been fully cleaned and verified.  
# All numeric columns (Sleep Duration, Quality of Sleep, Stress Level, Activity)  
# contain valid values within expected ranges.
```

In [28]:

```
# BSAN 360 - Week 07 Project  
# Student Lucas Doyle  
# Topics String Processing, Data Wrangling, Reshaping & Pivoting
```

In [29]:

```
# String Processing  
sleep.dtypes
```

Out[29]:

```
Person ID          int64  
Gender            object  
Age               int64  
Occupation        object  
Sleep Duration    float64  
Quality of Sleep   int64  
Activity           int64  
Stress Level      int64  
BMI Category      object  
Blood Pressure    object  
Heart Rate         int64  
Daily Steps        int64  
Sleep Disorder    object  
dtype: object
```

In [30]:

```
# String Processing  
string_cols = ['Gender', 'Occupation', 'BMI Category', 'Blood Pressure', 'Sleep Dis  
  
for col in string_cols:  
    print(f"\nColumn: {col}")  
    print(sleep[col].unique()[:10])
```

```
Column: Gender
['Male' 'Female']

Column: Occupation
['Software Engineer' 'Doctor' 'Sales Representative' 'Teacher' 'Nurse'
 'Engineer' 'Accountant' 'Scientist' 'Lawyer' 'Salesperson']

Column: BMI Category
['Overweight' 'Normal' 'Obese' 'Normal Weight']

Column: Blood Pressure
['126/83' '125/80' '140/90' '120/80' '132/87' '130/86' '117/76' '118/76'
 '128/85' '131/86']

Column: Sleep Disorder
[nan 'Sleep Apnea' 'Insomnia']
```

```
In [31]: # Processing Strings

string_cols = ['Gender', 'Occupation', 'BMI Category', 'Blood Pressure', 'Sleep Dis
sleep[string_cols] = (sleep[string_cols]
    .apply(lambda s: s.str.strip().str.replace(r'\s+', ' ', regex=True)
    .apply(lambda s: s.str.title())))

sleep[string_cols].head()
```

```
Out[31]:
```

	Gender	Occupation	BMI Category	Blood Pressure	Sleep Disorder
0	Male	Software Engineer	Overweight	126/83	NaN
1	Male	Doctor	Normal	125/80	NaN
2	Male	Doctor	Normal	125/80	NaN
3	Male	Sales Representative	Obese	140/90	Sleep Apnea
4	Male	Sales Representative	Obese	140/90	Sleep Apnea

```
In [32]: # Processing Strings

sleep['BMI Category'] = sleep['BMI Category'].replace({'Normal Weight': 'Normal'})

sleep['BMI Category'].value_counts()
```

```
Out[32]:
```

```
BMI Category
Normal      216
Overweight   148
Obese       10
Name: count, dtype: int64
```

```
In [33]: # Processing Strings

bp = sleep['Blood Pressure'].str.strip()
sleep[['Systolic', 'Diastolic']] = bp.str.split('/', expand=True).astype('int64')
```

```
sleep[['Blood Pressure', 'Systolic', 'Diastolic']].head()
```

Out[33]:

	Blood Pressure	Systolic	Diastolic
0	126/83	126	83
1	125/80	125	80
2	125/80	125	80
3	140/90	140	90
4	140/90	140	90

In [34]:

```
# Processing Strings  
sleep[['Gender', 'Occupation', 'BMI Category', 'Sleep Disorder']].head()
```

Out[34]:

	Gender	Occupation	BMI Category	Sleep Disorder
0	Male	Software Engineer	Overweight	NaN
1	Male	Doctor	Normal	NaN
2	Male	Doctor	Normal	NaN
3	Male	Sales Representative	Obese	Sleep Apnea
4	Male	Sales Representative	Obese	Sleep Apnea

In [35]:

```
# Part 2: Data Wrangling - Combining and Merging Datasets  
import pandas as pd  
  
datasets = [name for name in globals() if isinstance(globals()[name], pd.DataFrame)]  
print("Available DataFrames:", datasets)
```

Available DataFrames: ['__', '_4', 'sleep', '_7', '_8', '_11', '_12', '_15', '_18', 'sleep_project', 'stress_sleep', 'activity_sleep', '_21', '_26', '_31', '_33', '_34']

In [36]:

```
# Part 2: Data Wrangling - Combining and Merging Datasets  
# Since only one primary dataset ('sleep') is used in this project, merging or conc  
# The smaller DataFrames ('sleep_project', 'stress_sleep', 'activity_sleep') are su  
# created for specific analyses, not separate files. Therefore, no merge operations
```

In [37]:

```
# Part 3 - Data Wrangling: Reshaping and Pivoting  
  
sleep.info()  
sleep.head()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 374 entries, 0 to 373
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Person ID        374 non-null    int64  
 1   Gender           374 non-null    object  
 2   Age              374 non-null    int64  
 3   Occupation       374 non-null    object  
 4   Sleep Duration   374 non-null    float64 
 5   Quality of Sleep 374 non-null    int64  
 6   Activity          374 non-null    int64  
 7   Stress Level     374 non-null    int64  
 8   BMI Category     374 non-null    object  
 9   Blood Pressure   374 non-null    object  
 10  Heart Rate        374 non-null    int64  
 11  Daily Steps       374 non-null    int64  
 12  Sleep Disorder   155 non-null    object  
 13  Systolic          374 non-null    int64  
 14  Diastolic         374 non-null    int64  
dtypes: float64(1), int64(9), object(5)
memory usage: 44.0+ KB

```

Out[37]:

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Activity	Stress Level	BMI Category	Pr
0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight	
1	2	Male	28	Doctor	6.2	6	60	8	Normal	
2	3	Male	28	Doctor	6.2	6	60	8	Normal	
3	4	Male	28	Sales Representative	5.9	4	30	8	Obese	
4	5	Male	28	Sales Representative	5.9	4	30	8	Obese	



In [38]: # Part 3: Reshaping & Pivoting

```

# Mean sleep metrics by Gender x BMI Category
df_pivot1 = (sleep
              .pivot_table(index='Gender',
                           columns='BMI Category',
                           values=['Sleep Duration','Quality of Sleep'],
                           aggfunc='mean')
              .round(2))
df_pivot1

```

Out[38]:

BMI Category	Quality of Sleep			Sleep Duration		
	Normal	Obese	Overweight	Normal	Obese	Overweight
Gender						
Female	8.28	7.00	7.22	7.70	7.40	6.88
Male	7.28	6.33	6.10	7.21	6.91	6.49

In [39]:

```
# Part 3: Reshaping & Pivoting

# Mean sleep/stress/activity by Occupation
df_pivot2 = (sleep
              .pivot_table(index='Occupation',
                           values=['Sleep Duration','Quality of Sleep','Stress Level',
                                   aggfunc='mean')
              .round(2))
df_pivot2
```

Out[39]:

Occupation	Activity	Quality of Sleep	Sleep Duration	Stress Level
	Accountant	7.89	7.11	4.59
Doctor	55.35	6.65	6.97	6.73
Engineer	51.86	8.41	7.99	3.89
Lawyer	70.43	7.89	7.41	5.06
Manager	55.00	7.00	6.90	5.00
Nurse	78.59	7.37	7.06	5.55
Sales Representative	30.00	4.00	5.90	8.00
Salesperson	45.00	6.00	6.40	7.00
Scientist	41.00	5.00	6.00	7.00
Software Engineer	48.00	6.50	6.75	6.00
Teacher	45.62	6.98	6.69	4.53

In [40]:

```
# Part 3: Reshaping & Pivoting
# Melt to Long format for plotting/analysis

metrics = ['Sleep Duration','Quality of Sleep','Stress Level','Activity']
df_long = (sleep[['Person ID']] + metrics)
           .melt(id_vars='Person ID', var_name='Metric', value_name='Value')
df_long.head()
```

Out[40]:

	Person ID	Metric	Value
0	1	Sleep Duration	6.1
1	2	Sleep Duration	6.2
2	3	Sleep Duration	6.2
3	4	Sleep Duration	5.9
4	5	Sleep Duration	5.9

In [41]:

```
# Part 3: Reshaping & Pivoting
# Pivot back to wide
df_wide = df_long.pivot(index='Person ID', columns='Metric', values='Value')
df_wide.head()
```

Out[41]:

	Metric	Activity	Quality of Sleep	Sleep Duration	Stress Level
Person ID					
1	42.0		6.0	6.1	6.0
2	60.0		6.0	6.2	8.0
3	60.0		6.0	6.2	8.0
4	30.0		4.0	5.9	8.0
5	30.0		4.0	5.9	8.0

In [42]:

```
# Project 4
# For the data you have collected for your project, perform the following task(s):

# Plotting and Visualization
# A few things to watch out for:

# Any visualization you do for this week is expected to be geared towards exploratory analysis
# You can use the lecture notebook and/or the hands-on exercises to guide you with
```

In [43]:

```
#4
import matplotlib.pyplot as plt
```

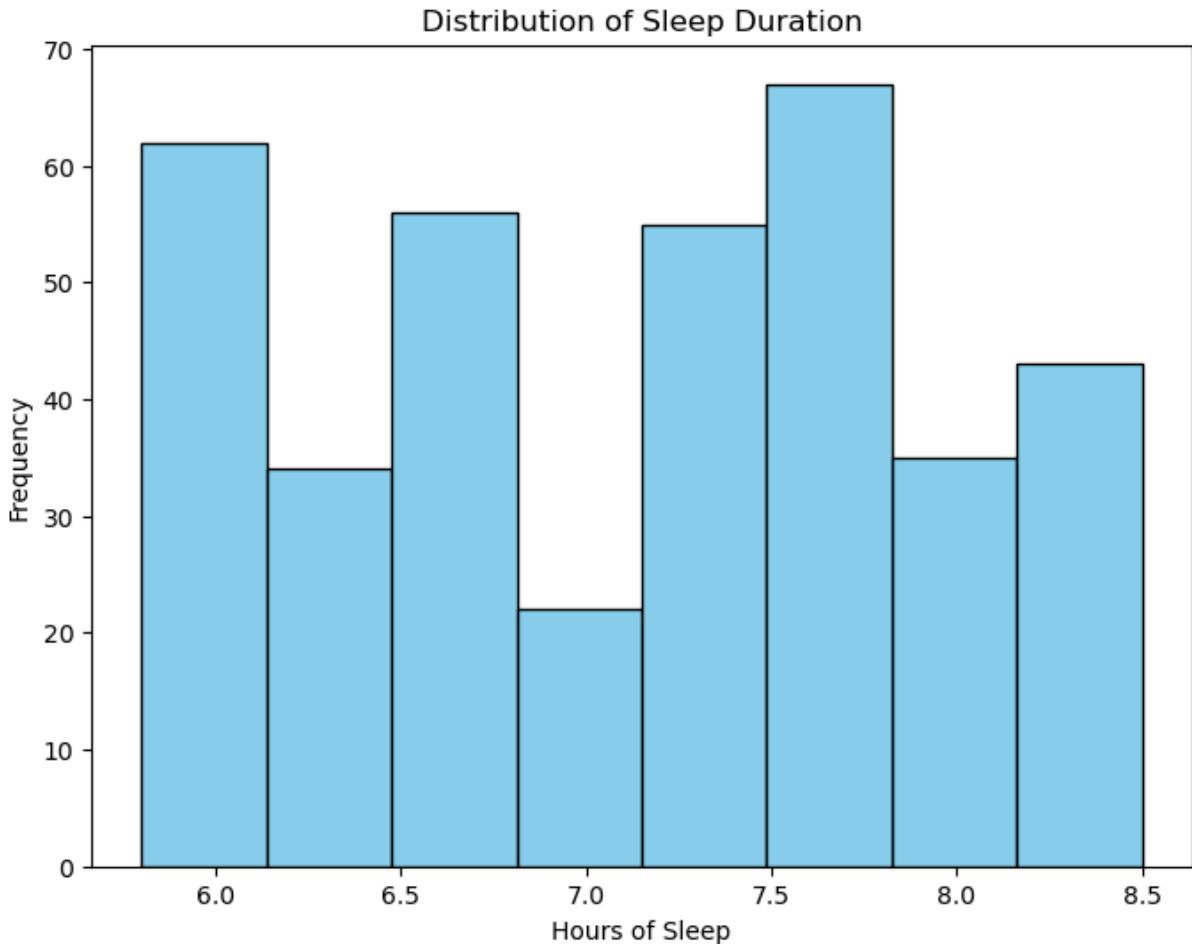
In [44]:

```
# Plotting and Visualization - histogram

plt.figure(figsize=(8,6))
plt.hist(sleep['Sleep Duration'], bins=8, color='skyblue', edgecolor='black')

plt.xlabel("Hours of Sleep")
plt.ylabel("Frequency")
plt.title("Distribution of Sleep Duration")

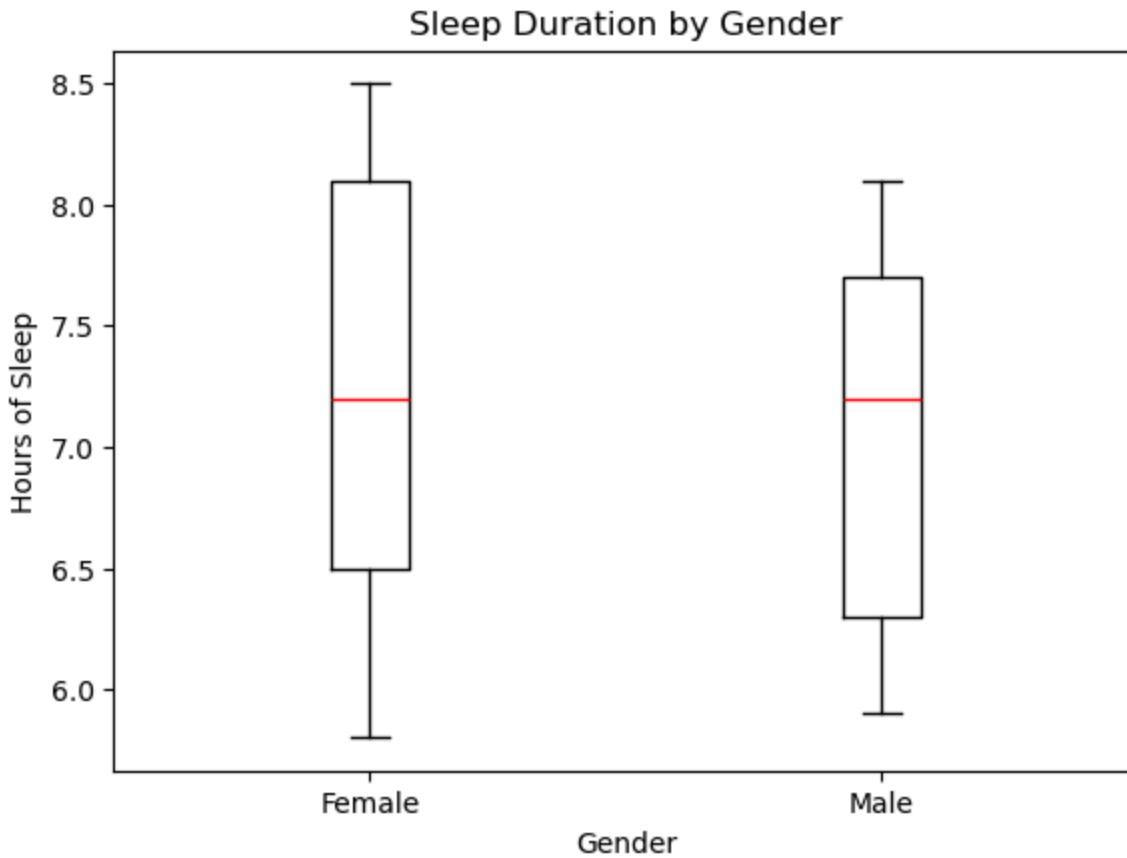
plt.show()
```



```
In [45]: # The histogram shows the distribution of sleep duration in my dataset.  
# It lets me see the most common sleep ranges and if the data is skewed or spread.  
# fairly even with no strong skew
```

```
In [46]: # 4 boxplot of sleep duration by gender  
  
plt.figure(figsize=(8,6))  
sleep.boxplot(column='Sleep Duration', by='Gender', grid=False,  
              boxprops=dict(color='black'),  
              medianprops=dict(color='red'),  
              whiskerprops=dict(color='black'),  
              capprops=dict(color='black'))  
  
plt.xlabel("Gender")  
plt.ylabel("Hours of Sleep")  
plt.title("Sleep Duration by Gender")  
plt.suptitle("") # remove default title  
  
plt.show()
```

<Figure size 800x600 with 0 Axes>

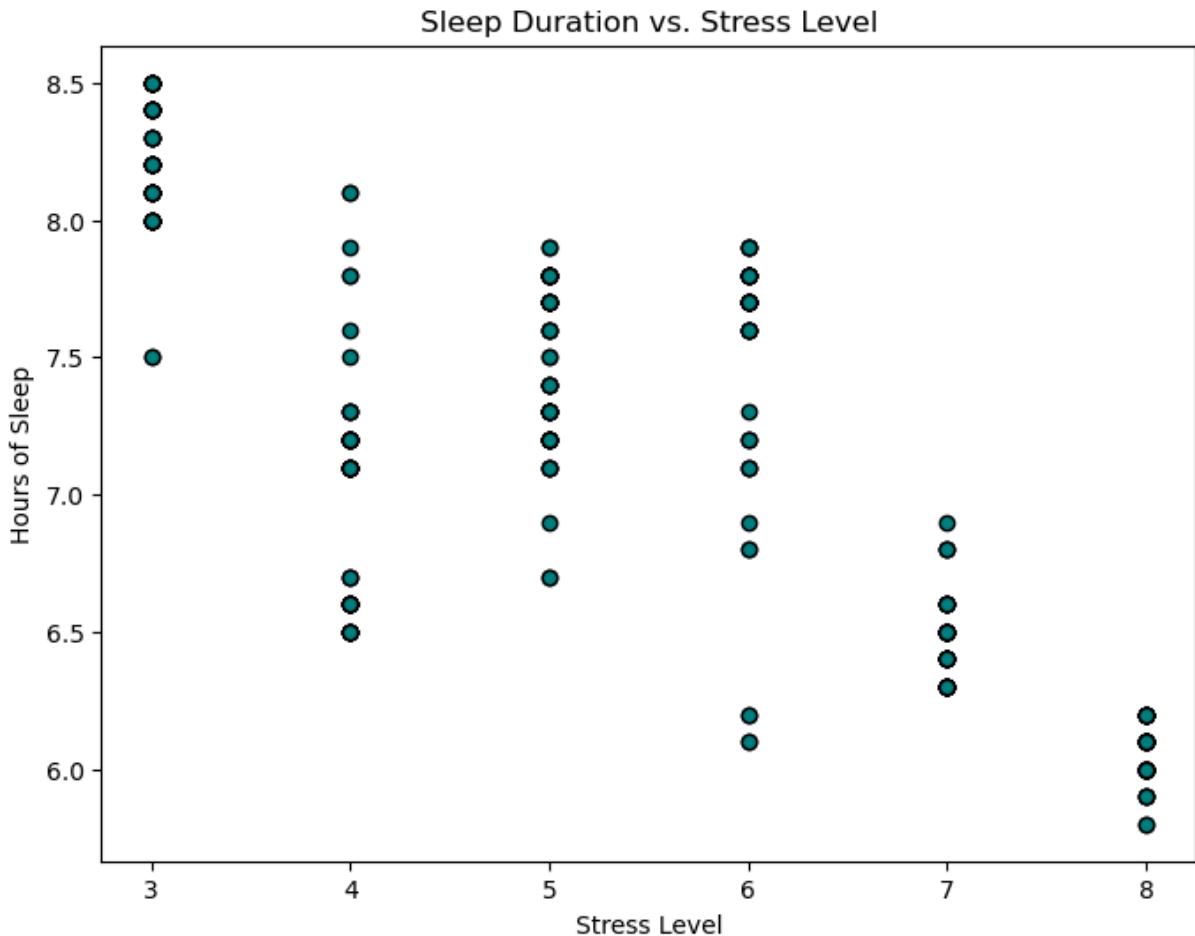


```
In [47]: # Sleep duration looks very similar between males and females, with both groups cen
```

```
In [48]: # scatterplot of sleep duration vs stress level
plt.figure(figsize=(8,6))
plt.scatter(sleep['Stress Level'], sleep['Sleep Duration'],
            color='teal', edgecolors='black')

plt.xlabel("Stress Level")
plt.ylabel("Hours of Sleep")
plt.title("Sleep Duration vs. Stress Level")

plt.show()
```



```
In [49]: # As stress level increases, sleep duration tends to decrease slightly, showing a linear relationship.
```

```
In [50]: # 4 boxplot of sleep duration by sleep disorder
```

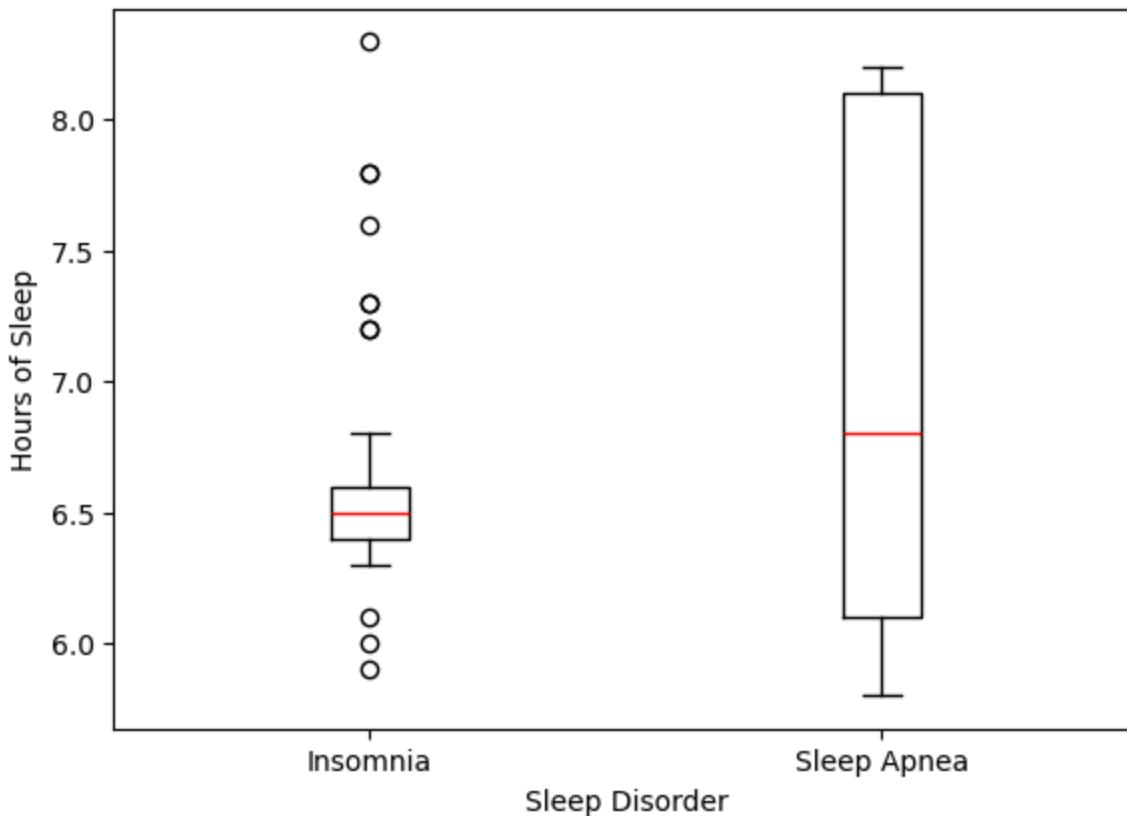
```
plt.figure(figsize=(8,6))
sleep.boxplot(column='Sleep Duration', by='Sleep Disorder', grid=False,
              boxprops=dict(color='black'),
              medianprops=dict(color='red'),
              whiskerprops=dict(color='black'),
              capprops=dict(color='black'))

plt.xlabel("Sleep Disorder")
plt.ylabel("Hours of Sleep")
plt.title("Sleep Duration by Sleep Disorder Status")
plt.suptitle("") # remove default title

plt.show()
```

<Figure size 800x600 with 0 Axes>

Sleep Duration by Sleep Disorder Status



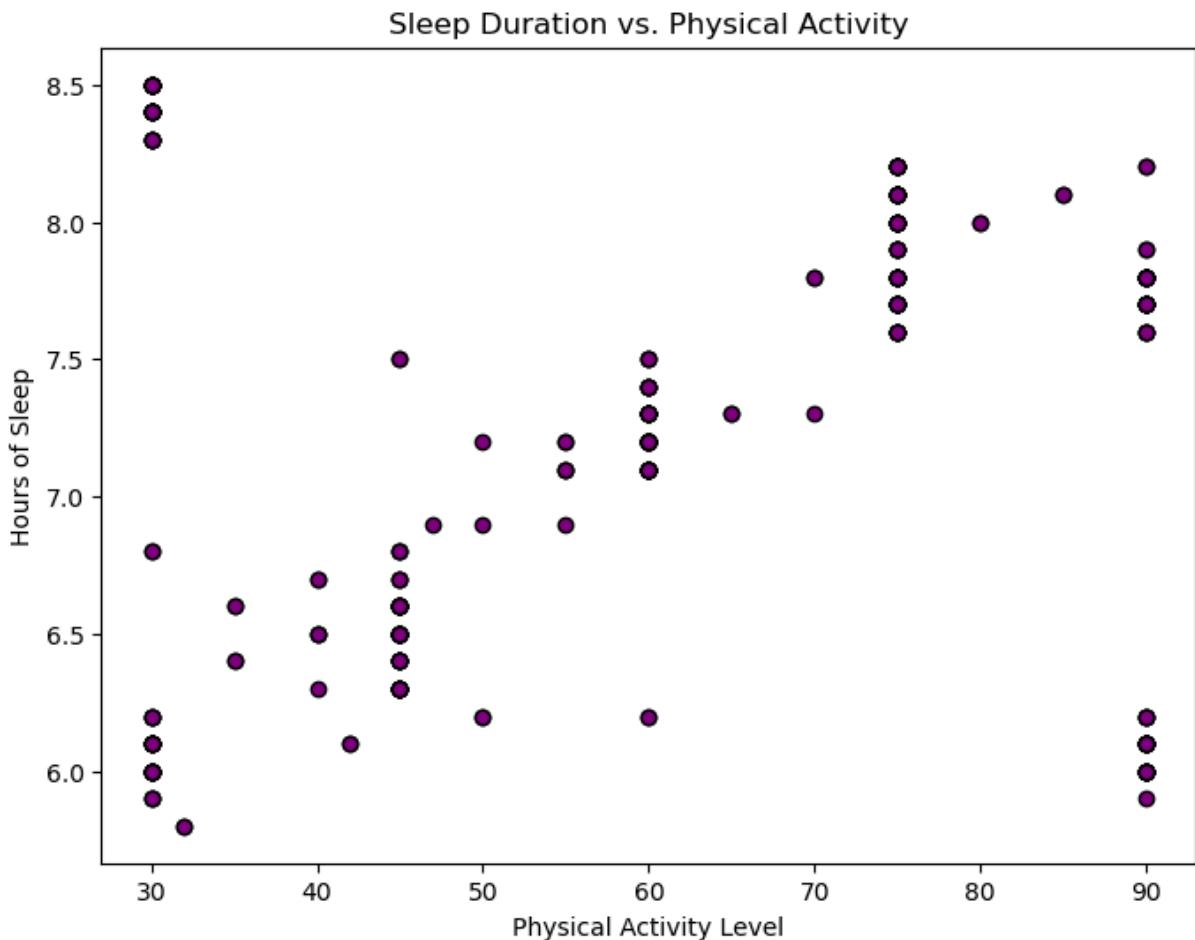
```
In [51]: # People with insomnia tend to sleep less on average than those with sleep apnea, so
```

```
In [52]: # 4 scatterplot of sleep duration vs physical activity
```

```
plt.figure(figsize=(8,6))
plt.scatter(sleep['Activity'], sleep['Sleep Duration'],
            color='purple', edgecolors='black')

plt.xlabel("Physical Activity Level")
plt.ylabel("Hours of Sleep")
plt.title("Sleep Duration vs. Physical Activity")

plt.show()
```



```
In [53]: # Higher physical activity Levels are generally associated with Longer sleep duration
```

```
In [54]: # Note
# Based on your feedback, the original scatterplots for Sleep Duration vs Stress Level
# and Sleep Duration vs Physical Activity Level created vertical streaks due to the
# x-axis containing only a few integer values.

#
# Instead of using the scatterplots above, I am now using bubble charts to reveal the
# true density at each combination of values. Each bubble represents a unique pair
# of the two variables, with the bubble size proportional to the number of overlapping
# data points. This removes the overplotting.
```

```
In [55]: # Note
sleep['Sleep_round'] = sleep['Sleep Duration'].round(1)

combo = (
    sleep.groupby(['Activity', 'Sleep_round'])
        .size()
        .reset_index(name='count')
)
```

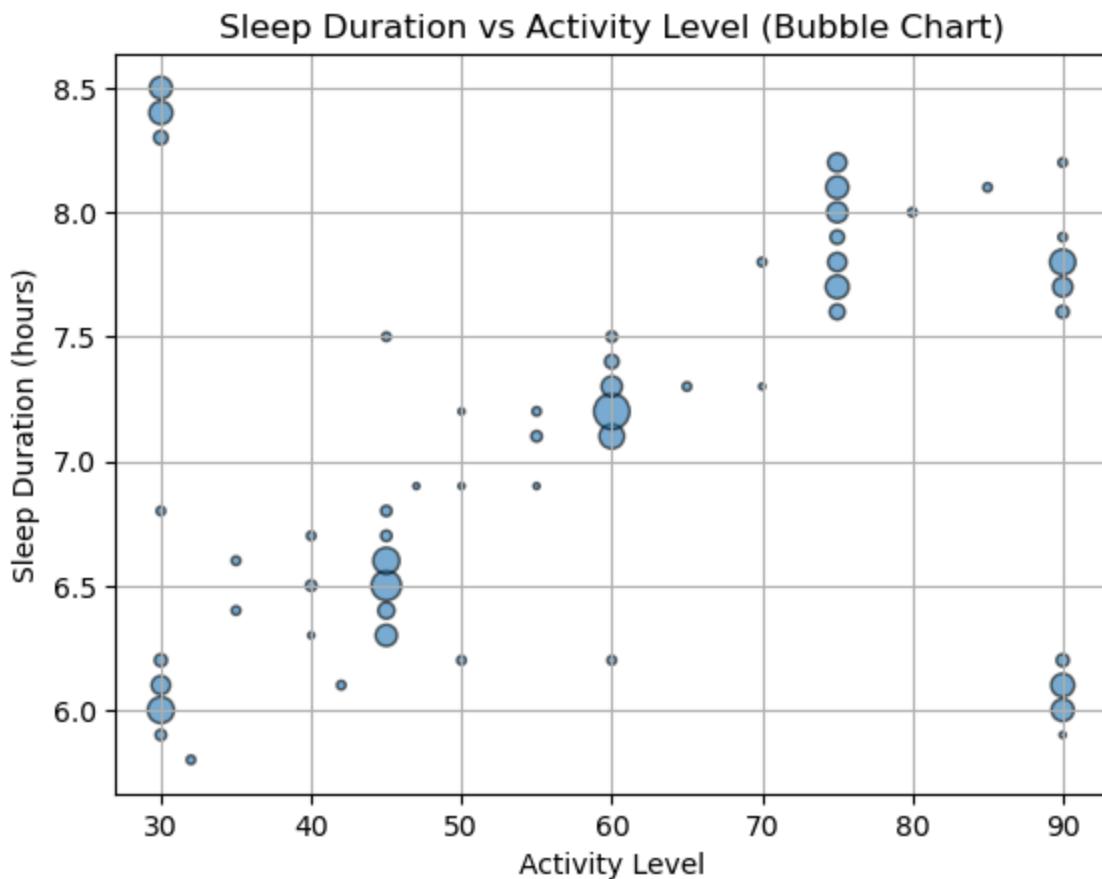
```
In [56]: # Note
plt.scatter(
    combo['Activity'],
    combo['Sleep_round'],
    s=combo['count'] * 5,  # Bubble size scaled
```

```

        alpha=0.6,
        edgecolors='black'
    )

plt.title('Sleep Duration vs Activity Level (Bubble Chart)')
plt.xlabel('Activity Level')
plt.ylabel('Sleep Duration (hours)')
plt.grid(True)
plt.show()

```



```

In [57]: # Note
sleep['Sleep_round'] = sleep['Sleep Duration'].round(1)

combo2 = (
    sleep.groupby(['Stress Level', 'Sleep_round'])
        .size()
        .reset_index(name='count')
)

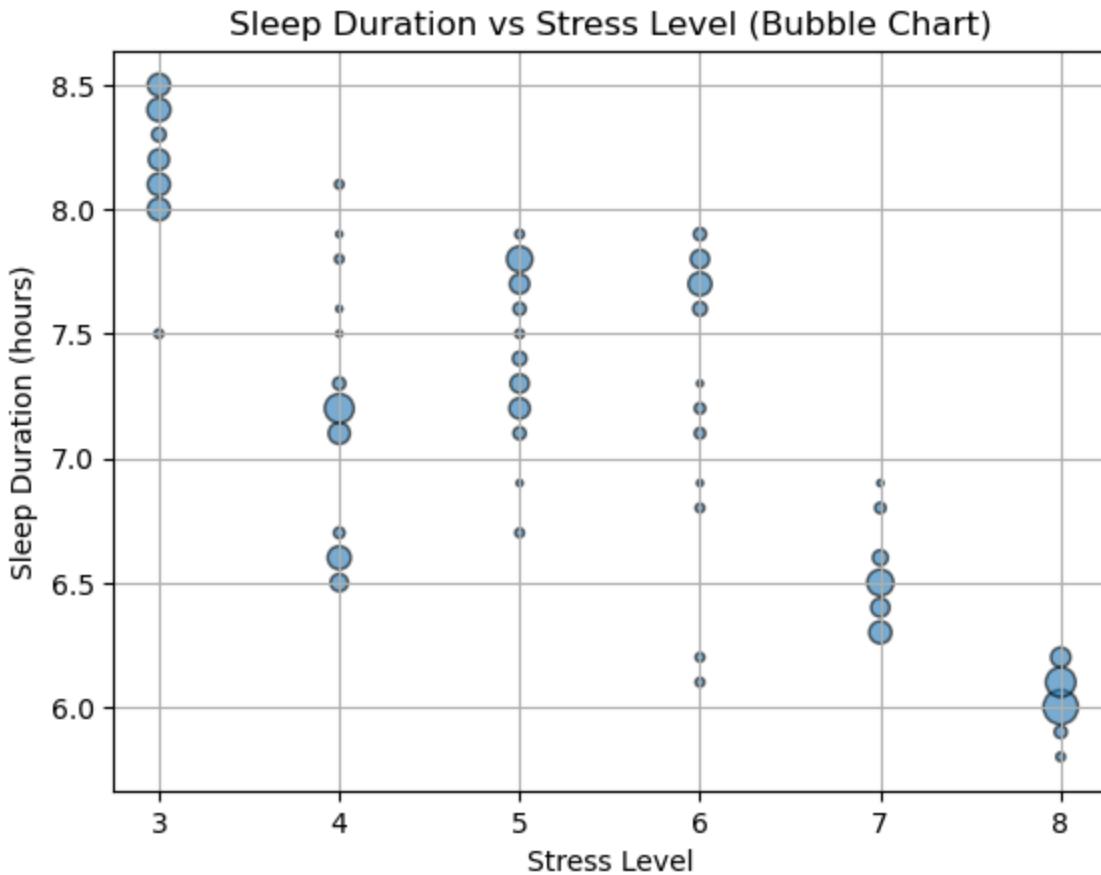
```

```

In [58]: # Note
plt.scatter(
    combo2['Stress Level'],
    combo2['Sleep_round'],
    s=combo2['count'] * 5,
    alpha=0.6,
    edgecolors='black'
)

```

```
plt.title('Sleep Duration vs Stress Level (Bubble Chart)')
plt.xlabel('Stress Level')
plt.ylabel('Sleep Duration (hours)')
plt.grid(True)
plt.show()
```



```
In [59]: # Project Assignment 5
# For the data you have collected for your project, perform the following task(s):

# Data Aggregation and Group Operations
# Grouping Data
# Pivot Tables and Cross-Tabulation
```

```
In [60]: # 1 Data Aggregation and Group Operations

sleep.describe()
```

Out[60]:

	Person ID	Age	Sleep Duration	Quality of Sleep	Activity	Stress Level	Heart Rate
count	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000
mean	187.500000	42.184492	7.132086	7.312834	59.171123	5.385027	70.165775
std	108.108742	8.673133	0.795657	1.196956	20.830804	1.774526	4.135676
min	1.000000	27.000000	5.800000	4.000000	30.000000	3.000000	65.000000
25%	94.250000	35.250000	6.400000	6.000000	45.000000	4.000000	68.000000
50%	187.500000	43.000000	7.200000	7.000000	60.000000	5.000000	70.000000
75%	280.750000	50.000000	7.800000	8.000000	75.000000	7.000000	72.000000
max	374.000000	59.000000	8.500000	9.000000	90.000000	8.000000	86.000000

In [61]:

```
# 1
sleep['Age_Group'] = pd.cut(
    sleep['Age'],
    bins=[25, 35, 45, 55, 65],
    labels=['26-35', '36-45', '46-55', '56-65']
).round(2)

sleep.groupby('Age_Group')[['Sleep Duration', 'Quality of Sleep', 'Stress Level']].mean()
```

C:\Users\lucas\AppData\Local\Temp\ipykernel_26952\839748768.py:8: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
sleep.groupby('Age_Group')[['Sleep Duration', 'Quality of Sleep', 'Stress Level']].mean()
```

Out[61]:

Age Group	Sleep Duration	Quality of Sleep	Stress Level
26-35	6.854255	6.500000	6.574468
36-45	7.050588	7.341176	5.135294
46-55	7.236364	7.519481	5.506494
56-65	8.100000	9.000000	3.000000

In [62]:

2 Grouping Data

```
sleep.groupby('Stress Level')[['Sleep Duration', 'Activity']].mean()
```

Out[62]:

Sleep Duration Activity

Stress Level	Sleep Duration	Activity
3	8.226761	54.718310
4	7.030000	55.785714
5	7.483582	74.253731
6	7.454348	67.152174
7	6.468000	43.840000
8	6.050000	58.342857

In [63]:

```
# 2
sleep['Activity_Group'] = pd.cut(
    sleep['Activity'],
    bins=[0, 40, 70, 100],
    labels=['Low', 'Medium', 'High']
).round(2)

# Group by Activity_Group
sleep.groupby('Activity_Group')[['Sleep Duration', 'Quality of Sleep', 'Stress Level']].mean()
```

C:\Users\lucas\AppData\Local\Temp\ipykernel_26952\1441615967.py:9: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

 sleep.groupby('Activity_Group')[['Sleep Duration', 'Quality of Sleep', 'Stress Level']].mean()

Out[63]:

Sleep Duration Quality of Sleep Stress Level

Activity_Group	Sleep Duration	Quality of Sleep	Stress Level
Low	7.067500	7.037500	5.775000
Medium	6.880128	7.211538	5.205128
High	7.454348	7.586957	5.362319

In [64]:

```
# 3 Pivot Tables and Cross-Tabulation

sleep.pivot_table(
    values='Sleep Duration',
    index='Age_Group',
    columns='Activity_Group',
    aggfunc='mean'
).round(2)
```

C:\Users\lucas\AppData\Local\Temp\ipykernel_26952\379069570.py:3: FutureWarning: The default value of observed=False is deprecated and will change to observed=True in a future version of pandas. Specify observed=False to silence this warning and retain the current behavior

 sleep.pivot_table(

```
Out[64]: Activity_Group  Low  Medium  High
```

Age_Group		Low	Medium	High
26–35	6.10	7.02	7.74	
36–45	6.66	6.87	7.76	
46–55	8.43	6.79	6.29	
56–65	NaN	NaN	8.10	

```
In [65]:
```

```
# 3
pd.crosstab(
    sleep['Stress Level'],
    sleep['Activity_Group']
)
```

```
Out[65]: Activity_Group  Low  Medium  High
```

Stress Level		Low	Medium	High
3	32	2	37	
4	0	66	4	
5	2	32	33	
6	2	12	32	
7	8	42	0	
8	36	2	32	

```
In [66]:
```

```
# Lab 6
# This week, we covered some advanced topics (Time Series Data and JSON Data) which
# What is relevant, however, is that you are now ready to learn and code advanced t
# For this week, you will need to finish any remaining analyses you need to conduct
# This may require some research on your part.
# For example, if you are looking to compare three groups to each other, you will n
# The statistics packages make most analysis tasks fairly straightforward, and plea
```

```
In [67]:
```

```
# Remaining analysis
# Correlation Heatmap
# Regression Model
# Regression Visualizations
# ANOVA
```

```
In [68]:
```

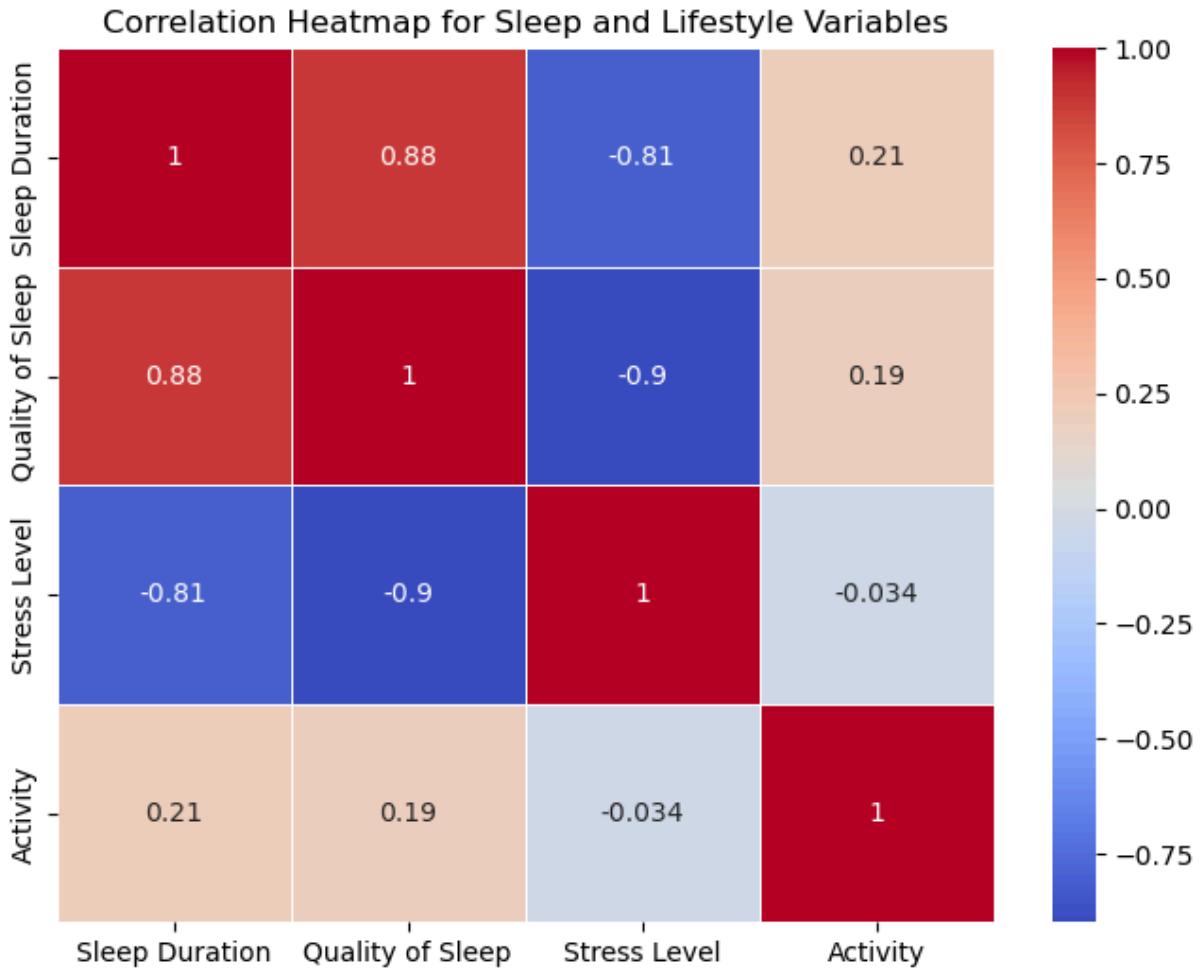
```
# Correlation Heatmap

import seaborn as sns

corr_vars = sleep[['Sleep Duration', 'Quality of Sleep', 'Stress Level', 'Activity']]
corr_matrix = corr_vars.corr()

plt.figure(figsize=(8, 6))
```

```
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Correlation Heatmap for Sleep and Lifestyle Variables')
plt.show()
```



```
In [69]: # Regression Model 1
# Whether lifestyle factors (stress level and physical activity) predict total sleep

import statsmodels.api as sm

X1 = sleep[['Stress Level', 'Activity']]
y1 = sleep['Sleep Duration']

X1 = sm.add_constant(X1)

model1 = sm.OLS(y1, X1).fit()

print("Model 1: Sleep Duration ~ Stress Level + Activity")
print(model1.summary())
```

Model 1: Sleep Duration ~ Stress Level + Activity

OLS Regression Results

Dep. Variable:	Sleep Duration	R-squared:	0.692
Model:	OLS	Adj. R-squared:	0.690
Method:	Least Squares	F-statistic:	416.6
Date:	Wed, 10 Dec 2025	Prob (F-statistic):	1.42e-95
Time:	20:36:17	Log-Likelihood:	-224.53
No. Observations:	374	AIC:	455.1
Df Residuals:	371	BIC:	466.8
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	8.6572	0.100	86.885	0.000	8.461	8.853
Stress Level	-0.3608	0.013	-27.908	0.000	-0.386	-0.335
Activity	0.0071	0.001	6.412	0.000	0.005	0.009

Omnibus:	3.064	Durbin-Watson:	0.656
Prob(Omnibus):	0.216	Jarque-Bera (JB):	3.134
Skew:	-0.202	Prob(JB):	0.209
Kurtosis:	2.807	Cond. No.	275.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [70]: # Reggesion 2
# Predicting Quality of Sleep
X2 = sleep[['Sleep Duration', 'Stress Level', 'Activity']]
y2 = sleep['Quality of Sleep']

X2 = sm.add_constant(X2)

model2 = sm.OLS(y2, X2).fit()

print("\nModel 2: Quality of Sleep ~ Sleep Duration + Stress Level + Activity")
print(model2.summary())
```

Model 2: Quality of Sleep ~ Sleep Duration + Stress Level + Activity
OLS Regression Results

Dep. Variable:	Quality of Sleep	R-squared:	0.884
Model:	OLS	Adj. R-squared:	0.883
Method:	Least Squares	F-statistic:	942.3
Date:	Wed, 10 Dec 2025	Prob (F-statistic):	8.05e-173
Time:	20:36:17	Log-Likelihood:	-194.18
No. Observations:	374	AIC:	396.4
Df Residuals:	370	BIC:	412.0
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	4.7503	0.425	11.176	0.000	3.914	5.586
Sleep Duration	0.6070	0.048	12.663	0.000	0.513	0.701
Stress Level	-0.3835	0.021	-18.248	0.000	-0.425	-0.342
Activity	0.0050	0.001	4.708	0.000	0.003	0.007

Omnibus:	52.471	Durbin-Watson:	1.060
Prob(Omnibus):	0.000	Jarque-Bera (JB):	81.235
Skew:	-0.876	Prob(JB):	2.29e-18
Kurtosis:	4.463	Cond. No.	1.28e+03

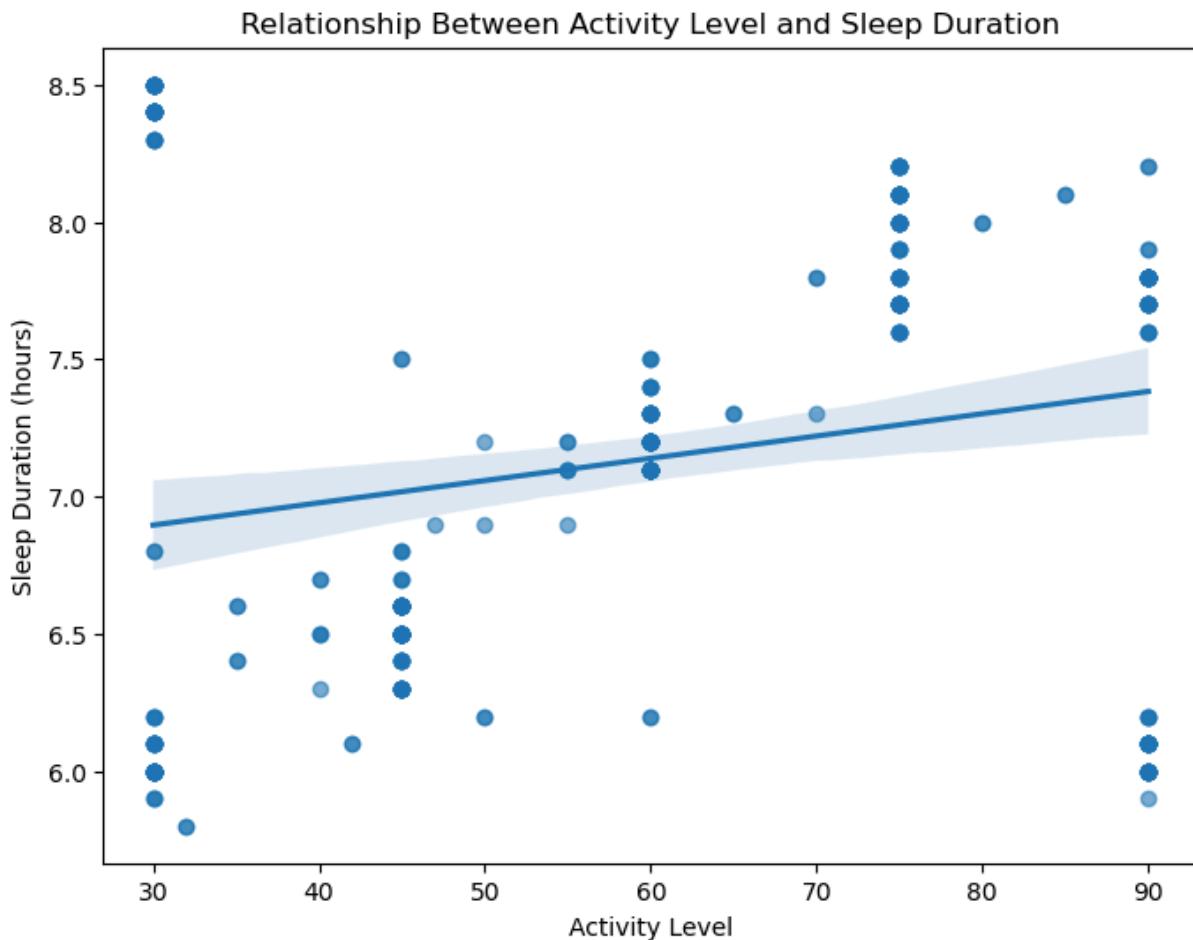
Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.28e+03. This might indicate that there are strong multicollinearity or other numerical problems.

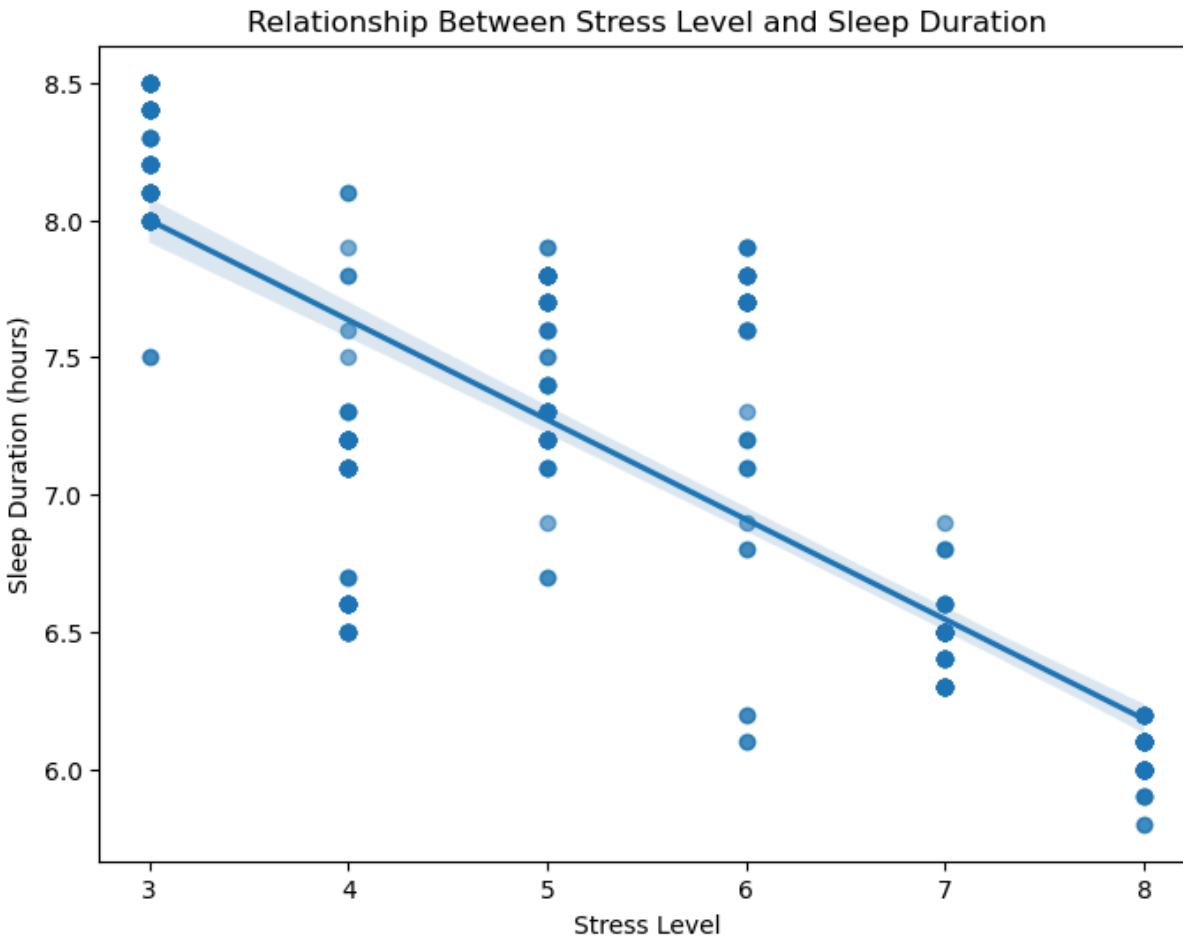
In [71]:

```
# Regression Visualizations
# how Activity and Stress Level relate to Sleep Duration.

plt.figure(figsize=(8, 6))
sns.regplot(data=sleep, x='Activity', y='Sleep Duration', scatter_kws={'alpha': 0.6}
plt.title('Relationship Between Activity Level and Sleep Duration')
plt.xlabel('Activity Level')
plt.ylabel('Sleep Duration (hours)')
plt.show()
```



```
In [72]: # Regression Visualizations 2
# Sleep Duration vs Stress Level
plt.figure(figsize=(8, 6))
sns.regplot(data=sleep, x='Stress Level', y='Sleep Duration', scatter_kws={'alpha': 0.5})
plt.title('Relationship Between Stress Level and Sleep Duration')
plt.xlabel('Stress Level')
plt.ylabel('Sleep Duration (hours)')
plt.show()
```



```
In [73]: # ANOVA - Sleep Duration across Activity Groups
import statsmodels.api as sm
import statsmodels.formula.api as smf
from statsmodels.stats.anova import anova_lm
```

```
In [74]: # ANOVA
anova_model = smf.ols('Q("Sleep Duration") ~ C(Activity_Group)', data=sleep).fit()
anova_results = anova_lm(anova_model, typ=2)

anova_results
```

```
Out[74]:
```

	sum_sq	df	F	PR(>F)
C(Activity_Group)	24.568685	2.0	21.541669	1.409553e-09
Residual	211.566289	371.0		NaN

```
In [75]: # Tukey test recommendation
from statsmodels.stats.multicomp import pairwise_tukeyhsd

tukey_results = pairwise_tukeyhsd(
    endog=sleep['Sleep Duration'],
    groups=sleep['Activity_Group'],
    alpha=0.05
)
```

```
print(tukey_results)
```

```
Multiple Comparison of Means - Tukey HSD, FWER=0.05
=====
group1 group2 meandiff p-adj    lower   upper  reject
-----
High    Low    -0.3868  0.0009 -0.6366 -0.1371  True
High   Medium  -0.5742    0.0 -0.7819 -0.3666  True
Low    Medium  -0.1874  0.1695 -0.4317   0.057 False
-----
```